



After lots of testing with various amounts of threads and depths, it is readily apparent that the scalability of the rlogin cluster contributes greatly to the speedup shown above. The more processors we have running the quicksort, the greater the speedup. The improvement rate starts to decline once you give it a certain amount of processors. The rate stops increasing fairly quickly, and drops off quicker and quicker the more processors you have. Eventually adding more processors won't help performance.

#Threads      Optimal Depth

2	12
4	14
6	16
8	16
12	18
14	17
16	18